

# Getting Started with SLURM and Python Environments on an HPC Cluster

Md Shahriar Forhad  
[shahriar.forhad.eee@gmail.com](mailto:shahriar.forhad.eee@gmail.com)  
<https://shahriar88.github.io/>

December 8, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Logging In to CRADLE . . . . .	3
<b>2</b>	<b>A Minimal SLURM Job Script Example</b>	<b>3</b>
<b>3</b>	<b>Understanding Key #SBATCH Directives</b>	<b>4</b>
3.1	Number of Nodes: -N . . . . .	4
3.2	Partition (Queue): -p . . . . .	5
3.3	GPU Request: -gres . . . . .	5
3.4	Time Limit: -time . . . . .	5
3.5	Summary Table of Core Directives . . . . .	5
<b>4</b>	<b>Why Use Python Environments on an HPC Cluster?</b>	<b>5</b>
<b>5</b>	<b>Cloning a Shared TensorFlow Environment</b>	<b>6</b>
5.1	Step 1: Install the cloning tool (one time) . . . . .	6
5.2	Step 2: Clone the shared environment . . . . .	6
5.3	Step 3: Activate the cloned environment . . . . .	7
5.4	Step 4: Add new packages (example: pandas) . . . . .	7
5.5	Step 5: Use the cloned environment in your SLURM job . . . . .	7
5.6	Mental Model Summary . . . . .	8
<b>6</b>	<b>Creating Your Own Python Virtual Environment</b>	<b>8</b>
6.1	Option A (Recommended for CPU-Only): Create from Scratch . . . . .	8
6.1.1	Step 1: Choose a location . . . . .	8
6.1.2	Step 2: Create the virtual environment . . . . .	8
6.1.3	Step 3: Activate the environment . . . . .	9
6.1.4	Step 4: Upgrade pip (recommended) . . . . .	9
6.1.5	Step 5: Install packages . . . . .	9
6.1.6	Step 6: Use this environment in a SLURM job . . . . .	9
6.2	Option B (Best for GPU): Clone an Optimized Shared Environment . . . . .	9

<b>7</b>	<b>Verifying Your Environment and Common Pitfalls</b>	<b>10</b>
7.1	Verifying Python and GPU Visibility . . . . .	10
7.2	Common Mistakes to Avoid . . . . .	10
7.3	Quick Decision Guide . . . . .	10
<b>8</b>	<b>Using Micromamba and Portable Conda Environments on the HPC Cluster</b>	<b>10</b>
8.1	Exporting a Portable Conda Environment on Your PC . . . . .	11
8.2	Copying the YAML File to CRADLE . . . . .	11
8.3	Installing Micromamba in Your Home Directory . . . . .	11
8.4	Creating a Micromamba Environment from the YAML File . . . . .	12
8.5	Activating the Environment on CRADLE . . . . .	12
8.6	Using a Micromamba Environment in a SLURM Job . . . . .	13
8.7	Verifying GPU Support Inside the Environment . . . . .	13
8.8	Best Practices and Common Pitfalls for Micromamba on HPC . . . . .	13
<b>9</b>	<b>Worked Examples: Environments and SLURM Job Scripts</b>	<b>14</b>
9.1	Creating a New CPU-Only Environment . . . . .	15
9.2	Cloning a Shared GPU-Optimized TensorFlow Environment . . . . .	15
9.3	Selecting an Environment in a Job Script . . . . .	15
9.4	Installing Libraries After Environment Activation . . . . .	16
9.5	Checking and Verifying Resources . . . . .	16
9.5.1	Check Available Partitions and Nodes . . . . .	16
9.5.2	Check Your Jobs . . . . .	16
9.5.3	Inspect Resources Allocated by SLURM . . . . .	16
9.5.4	Check Resources Inside a Job Script . . . . .	16
9.5.5	Check GPU Allocation . . . . .	17
9.5.6	Check CPU and Memory Usage . . . . .	17
9.5.7	Check Resources Used by Completed Jobs . . . . .	17
9.5.8	Minimal Resource Debug Job . . . . .	17
9.6	SLURM Job Templates for Different Resource Requirements . . . . .	18
9.6.1	Single Node, CPU-Only Job . . . . .	18
9.6.2	Single Node, Single GPU Job . . . . .	19
9.6.3	Single Node, Multiple GPUs (e.g., 4 GPUs) . . . . .	19
9.6.4	Multiple Nodes, CPU-Only (e.g., MPI or Distributed CPU Job) . . . . .	20
9.6.5	Multiple Nodes, Multiple GPUs (Distributed GPU Training) . . . . .	20
9.7	Quick Mapping: Which Example to Use? . . . . .	21
<b>10</b>	<b>Summary</b>	<b>21</b>

# 1 Introduction

This short tutorial explains how to:

- Write a basic SLURM batch script for running a TensorFlow job.
- Understand key `#SBATCH` directives, including nodes, partitions, GPUs, and time limits.
- Use and extend a shared TensorFlow environment on the cluster.
- Create your own Python virtual environments for CPU and GPU workloads.

All examples are designed to be copied directly into a terminal or SLURM job script on a typical HPC system (e.g., a university cluster).

## 1.1 Logging In to CRADLE

Once you have received your user account information, you can log in to the CRADLE cluster from a terminal or command-line prompt using your University credentials.

### Step 1: Connect to the University Network (VPN)

If you are off campus, you must first connect to the University network using the approved VPN software. Install and start the VPN client, then log in with your University credentials. Once the VPN connection is active, you can reach the CRADLE login node.

### Step 2: SSH into the Cluster

To log in, use `ssh` with your assigned username. In the examples below, replace `username` with your actual cluster username.

```
1 ssh username@login.cradle.university.edu
```

You will be prompted for your password:

```
1 username@login.cradle.university.edu's password:
```

After you enter your password correctly, you should see a prompt indicating that you are now on the cluster login node, similar to:

```
1 [username@login001 ~]$
```

At this point you are logged into CRADLE and ready to work with SLURM, Python environments, and the other tools described in this tutorial.

## 2 A Minimal SLURM Job Script Example

Listing 1 shows a basic SLURM job script that runs a TensorFlow test script in a shared environment.

Listing 1: Basic SLURM job script with TensorFlow environment

```
1 #!/bin/bash
2
3 # Job identification
4 #SBATCH --job-name=myFirstJob
5
6 # Standard output and error files (%j = job ID)
7 #SBATCH --output=myFirstJob.out.%j
8 #SBATCH --error=myFirstJob.err.%j
9
10 # Resource requests
11 #SBATCH -N 1
12 #SBATCH -p kimg
13 #SBATCH --gres=gpu:1
14 #SBATCH --time=1:00:00
15
16 # Display GPU assigned by SLURM
17 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
18
19 # Activate TensorFlow virtual environment
20 echo "Activating TensorFlow-2.6.2 environment"
21 source /shared/tensorflow-2.6.2/tf_env/bin/activate
22
23 # Run TensorFlow test script
24 echo "Running testTF.py"
25 python3 ~/testTFForSlurm/testTF.py
26
27 # Deactivate environment
28 echo "Deactivating TensorFlow-2.6.2 environment"
29 deactivate
```

To submit this job to SLURM, save the script as `myFirstJob.slurm` (for example) and run:

```
1 sbatch myFirstJob.slurm
```

### 3 Understanding Key #SBATCH Directives

The `#SBATCH` lines at the top of the script tell SLURM what resources you want. They are not shell commands; they are scheduler directives.

#### 3.1 Number of Nodes: -N

```
1 #SBATCH -N 1
```

This requests **1 compute node**. A node is a physical machine that may contain many CPU cores, memory, and possibly GPUs.

More examples:

```
1 #SBATCH -N 1 # Single-node TensorFlow training
2 #SBATCH -N 2 # Distributed job across 2 nodes
3 #SBATCH -N 1-4 # SLURM may allocate anywhere from 1 to 4 nodes
```

### 3.2 Partition (Queue): -p

```
1 #SBATCH -p kimq
```

The partition selects the *group of nodes* (queue) your job will run on.

Typical partitions might be:

- `compute` → Standard CPU nodes
- `bigmem` → High-memory nodes
- `gpu` or `kimq` → GPU-enabled nodes

Examples:

```
1 #SBATCH -p kimq # Run on GPU nodes in the kimq partition
2 #SBATCH -p compute # Run on CPU-only nodes
```

### 3.3 GPU Request: -gres

```
1 #SBATCH --gres=gpu:1
```

This requests **1 GPU** using the generic resource (GRES) mechanism.

More examples:

```
1 #SBATCH --gres=gpu:1 # One GPU
2 #SBATCH --gres=gpu:2 # Two GPUs
3 #SBATCH --gres=gpu:a100:1 # One NVIDIA A100 GPU (if supported)
```

**Important: Partition and GPU must match.** If you request GPUs but choose a CPU-only partition (e.g., `-p compute` with `-gres=gpu:1`), SLURM cannot satisfy the request and the job will stay pending or fail.

### 3.4 Time Limit: -time

```
1 #SBATCH --time=1:00:00
```

This sets a **wall-clock time limit** of 1 hour for your job. When this time is reached, SLURM terminates the job.

More examples:

```
1 #SBATCH --time=00:30:00 # 30 minutes
2 #SBATCH --time=04:00:00 # 4 hours
3 #SBATCH --time=2-00:00:00 # 2 days
```

### 3.5 Summary Table of Core Directives

## 4 Why Use Python Environments on an HPC Cluster?

On HPC clusters, shared software environments (for example, a shared TensorFlow installation) are usually:

- Optimized for the cluster hardware (correct CUDA, cuDNN, drivers).

Directive	Controls	Example
-N	Number of nodes	-N 1
-p	Hardware pool	-p klmq
-gres	GPU count	gpu:1
-time	Runtime limit	1:00:00

Table 1: Core SLURM directives and examples.

- Read-only for users.
- Not modifiable by individual users.

You *cannot* safely run:

```
1 pip install pandas
```

inside a shared environment such as

```
1 /shared/tensorflow-2.6.2/tf_env/
```

because:

- It is shared by everyone.
- You do not have write permission.
- Changes could break other users' jobs.

**Solution:** Clone the shared environment into your home directory, then customize the copy.

## 5 Cloning a Shared TensorFlow Environment

### 5.1 Step 1: Install the cloning tool (one time)

```
1 pip3 install --user virtualenv-clone
```

This:

- Installs `virtualenv-clone` into your user space.
- Does not require admin/root access.
- Only needs to be run once.

### 5.2 Step 2: Clone the shared environment

```
1 virtualenv-clone /shared/tensorflow-2.6.2/tf_env ~/myEnvs/tf_env
```

Here:

- `/shared/tensorflow-2.6.2/tf_env/` is the original, optimized, read-only TensorFlow environment.
- `~/myEnvs/tf_env` is your private, writable copy.

What gets copied:

- Python version.
- TensorFlow build (CUDA, cuDNN matched to cluster GPUs).
- All installed dependencies.
- Environment configuration tuned for the cluster.

Advantages:

- Runs as fast as the shared environment.
- You can safely add or upgrade packages.
- No conflicts with other users.

### 5.3 Step 3: Activate the cloned environment

```
1 source ~/myEnvs/tf_env/bin/activate
```

After activation:

- `python` and `pip` point to your cloned environment.
- Any `pip install` affects only your copy.

The shell prompt typically changes to something like:

```
1 (tf_env) username@node \$
```

### 5.4 Step 4: Add new packages (example: pandas)

With the environment active:

```
1 pip3 install pandas
```

What happens:

- `pandas` is installed only into `~/myEnvs/tf_env`.
- TensorFlow and CUDA support remain intact.
- The customization is isolated and safe.

You can also add:

```
1 pip install scikit-learn matplotlib seaborn tqdm
```

### 5.5 Step 5: Use the cloned environment in your SLURM job

Update the activation lines in your job script.

Old (shared environment):

```
1 echo "Activating TensorFlow-2.6.2 environment"
2 source /shared/tensorflow-2.6.2/tf_env/bin/activate
```

New (your cloned environment):

```
1 echo "Activating my custom environment"
2 source ~/myEnvs/tf_env/bin/activate
```

This ensures:

- Your job sees TensorFlow and any extra packages (e.g., pandas).
- You avoid dependency errors at runtime.
- The behavior of your job is reproducible.

## 5.6 Mental Model Summary

Piece	Purpose
Shared environment	Fast, optimized, read-only
Cloned environment	Same speed, user-modifiable
<code>virtualenv-clone</code>	Copies entire environment
<code>pip install</code>	Safe only in cloned environment
Job script	Must activate your chosen environment

Table 2: Conceptual roles of shared and cloned environments.

## 6 Creating Your Own Python Virtual Environment

Sometimes you want a completely independent Python environment. There are two main options; here we focus on standard `venv` and the cloned approach.

### 6.1 Option A (Recommended for CPU-Only): Create from Scratch

#### 6.1.1 Step 1: Choose a location

Create a directory to hold your environments:

```
1 mkdir -p ~/myEnvs
```

#### 6.1.2 Step 2: Create the virtual environment

```
1 python3 -m venv ~/myEnvs/myenv
```

This creates something like:

```
1 ~/myEnvs/myenv/
2 |- bin/
3 |- lib/
4 `-- pyvenv.cfg
```

### 6.1.3 Step 3: Activate the environment

```
1 source ~/myEnvs/myenv/bin/activate
```

Your shell prompt changes to:

```
1 (myenv) username@login $
```

### 6.1.4 Step 4: Upgrade pip (recommended)

```
1 pip install --upgrade pip
```

### 6.1.5 Step 5: Install packages

Example:

```
1 pip install tensorflow pandas numpy matplotlib
```

Note: On clusters, installing TensorFlow this way may not include GPU support unless CUDA and drivers are correctly matched. For GPU workloads, the cloned shared environment (Option B) is often safer.

### 6.1.6 Step 6: Use this environment in a SLURM job

In your job script:

```
1 source ~/myEnvs/myenv/bin/activate  
2 python train.py
```

## 6.2 Option B (Best for GPU): Clone an Optimized Shared Environment

This is essentially the workflow described in Section 5.

Step 1: Install the clone tool:

```
1 pip install --user virtualenv-clone
```

Step 2: Clone the shared environment:

```
1 virtualenv-clone /shared/tensorflow-2.6.2/tf_env ~/myEnvs/tf_env
```

Step 3: Activate it:

```
1 source ~/myEnvs/tf_env/bin/activate
```

Step 4: Add packages:

```
1 pip install pandas scikit-learn seaborn
```

Step 5: Update your SLURM job to activate:

```
1 source ~/myEnvs/tf_env/bin/activate
```

This preserves CUDA/cuDNN compatibility and performance while allowing you to customize the environment.

## 7 Verifying Your Environment and Common Pitfalls

### 7.1 Verifying Python and GPU Visibility

After activating an environment, verify what `python` and `pip` you are using:

```
1 which python  
2 which pip
```

For TensorFlow GPU detection:

```
1 python -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

You should see at least one GPU listed when running on a GPU node with correct drivers and environment.

### 7.2 Common Mistakes to Avoid

- Installing packages without activating the environment first.
- Using `pip install -user` inside a virtual environment.
- Mixing Conda environments and `venv` in the same workflow.
- Using the system Python (no virtual environment) for SLURM jobs.
- Requesting GPUs in a CPU-only partition.

### 7.3 Quick Decision Guide

Situation	Best choice
CPU-only work	Option A (new <code>venv</code> )
GPU + TensorFlow	Option B (clone shared env)
Cluster-optimized builds	Option B
Full custom control	Option A

Table 3: Choosing between creating and cloning environments.

## 8 Using Micromamba and Portable Conda Environments on the HPC Cluster

On many HPC clusters, including CRADLE, traditional Conda/Anaconda installations are discouraged or restricted because they are large, modify shell startup files aggressively, and may conflict with system libraries or CUDA drivers. A lightweight and HPC-friendly alternative is *Micromamba*, a single static binary that implements the Conda ecosystem without requiring a global base environment.

Micromamba allows you to:

- Recreate a Conda environment from a portable YAML file exported on your laptop.
- Install packages into your home directory without administrator rights.

- Activate environments explicitly inside SLURM job scripts.

The overall workflow is:

1. Export a portable Conda environment YAML on your local machine.
2. Copy the YAML file to CRADLE.
3. Install Micromamba in your home directory on CRADLE.
4. Create a Micromamba environment from the YAML file.
5. Activate the environment in your SLURM job scripts.

## 8.1 Exporting a Portable Conda Environment on Your PC

On your local machine where you already use Conda, start by creating a portable YAML description of your environment. Suppose your local environment is called `IR`.

First, activate it:

```
1 conda activate IR
```

Then export a *portable* environment file using `fromhistory`:

```
1 conda env export --from-history > IR_portable.yml
```

This creates a file `IR_portable.yml` containing the channels and high-level dependencies (without build hashes). It is much more likely to recreate successfully on a different system such as CRADLE.

## 8.2 Copying the YAML File to CRADLE

Use `scp` or an equivalent file transfer method to copy `IR_portable.yml` to your home directory on CRADLE. For example, from your local terminal:

```
1 scp IR_portable.yml username@login.cradle.university.edu:~
```

After logging into CRADLE, you should see the file in your home directory:

```
1 ls ~/IR_portable.yml
```

## 8.3 Installing Micromamba in Your Home Directory

Micromamba is distributed as a single binary, which makes it ideal for installation in a user home directory (no root access required).

Create a directory for the binary:

```
1 mkdir -p ~/bin
2 cd ~/bin
```

Download and unpack the latest Micromamba binary (Linux 64-bit):

```
1 curl -L https://micro.mamba.pm/api/micromamba/linux-64/latest \
2   | tar -xvj bin/micromamba
```

Add \$HOME/bin to your PATH in `~/.bashrc`:

```
1 echo 'export PATH=$HOME/bin:$PATH' >> ~/.bashrc
2 source ~/.bashrc
```

Verify that `micromamba` is available:

```
1 micromamba --version
```

## 8.4 Creating a Micromamba Environment from the YAML File

Choose a directory where Micromamba will store its environments, for example:

```
1 mkdir -p ~/micromamba/envs
```

Initialize Micromamba for Bash (one time):

```
1 micromamba shell init -s bash -p ~/micromamba
2 source ~/.bashrc
```

Now create a new environment (for example, `IR_NEW`) from your portable YAML file:

```
1 micromamba create \
2   -f ~/IR_portable.yml \
3   -n IR_NEW \
4   -p ~/micromamba/envs/IR_NEW
```

Here:

- `~/IR_portable.yml` is the YAML exported on your PC.
- `IR_NEW` is the environment name.
- `~/micromamba/envs/IR_NEW` is the full path Micromamba uses.

## 8.5 Activating the Environment on CRADLE

Once created, you can activate the environment on the CRADLE login node:

```
1 micromamba activate ~/micromamba/envs/IR_NEW
```

After activation:

- `python` and `pip` refer to the Micromamba environment.
- Any installs via `pip` or `micromamba install` affect only that environment.

You can verify:

```
1 which python
2 python --version
3 pip list
```

## 8.6 Using a Micromamba Environment in a SLURM Job

In a SLURM job script, you should explicitly set the Micromamba root and activate the environment before running Python code. For example:

```
1 #!/bin/bash
2 #SBATCH --job-name=ir_gpu_job
3 #SBATCH --output=ir_gpu_job.out.%j
4 #SBATCH --error=ir_gpu_job.err.%j
5 #SBATCH -N 1
6 #SBATCH -p kimg # GPU partition
7 #SBATCH --gres=gpu:1 # 1 GPU
8 #SBATCH --time=02:00:00 # 2 hours
9
10 # Ensure Micromamba root is defined
11 export MAMBA_ROOT_PREFIX=$HOME/micromamba
12
13 # Load shell initialization (if micromamba shell init was used)
14 source $HOME/.bashrc
15
16 echo "Host: $(hostname)"
17 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
18
19 echo "Activating Micromamba environment IR_NEW"
20 micromamba activate $HOME/micromamba/envs/IR_NEW
21
22 echo "Checking TensorFlow GPU visibility"
23 python -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
24
25 echo "Running training script"
26 python ~/projects/train_ir_model.py
```

This pattern ensures that:

- The correct Micromamba environment is active inside the job.
- GPU visibility can be checked via TensorFlow (or PyTorch).
- The job is reproducible and does not depend on the login-node Conda setup.

## 8.7 Verifying GPU Support Inside the Environment

To confirm that your Micromamba environment has working GPU support for TensorFlow (or another framework), you can run:

```
1 python -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

On a GPU node with correct drivers and CUDA libraries, this should list one or more GPU devices. You can also check the low-level GPU visibility with:

```
1 nvidia-smi
```

## 8.8 Best Practices and Common Pitfalls for Micromamba on HPC

**Best practices:**

- Export environments from Conda using `-from-history` or `-no-builds` to avoid hardware-specific build strings.
- Use Micromamba to recreate the environment on the cluster, not full Anaconda.
- Keep all Micromamba files under `~/micromamba` or a dedicated directory in your home space.
- Activate environments explicitly inside SLURM job scripts.

#### **Common pitfalls:**

- Installing full Anaconda on an HPC login node (disk usage and conflicts).
- Relying on Conda's `base` environment in batch jobs.
- Mixing Conda and Micromamba environments in the same workflow.
- Installing GPU frameworks with incompatible `cudatoolkit` versions; always check which CUDA versions the cluster supports.

Using Micromamba with a portable Conda YAML exported from your local machine gives you a reproducible, HPC-safe way to manage Python environments on CRADLE while respecting cluster policies and maximizing compatibility with the installed NVIDIA drivers and CUDA libraries.

## **9 Worked Examples: Environments and SLURM Job Scripts**

This section provides concrete, copy-pasteable examples covering:

1. Creating a new environment.
2. Cloning a shared TensorFlow environment.
3. Selecting (activating) an environment.
4. Installing extra libraries.
5. Checking and verifying resources.
6. Writing SLURM scripts for different resource requirements.

## 9.1 Creating a New CPU-Only Environment

```
1 # 1. Create a directory to store your environments
2 mkdir -p ~/myEnvs
3
4 # 2. Create a new CPU-only Python virtual environment
5 python3 -m venv ~/myEnvs/cpu_env
6
7 # 3. Activate the environment
8 source ~/myEnvs/cpu_env/bin/activate
9
10 # 4. Upgrade pip
11 pip install --upgrade pip
12
13 # 5. Install some common libraries
14 pip install numpy scipy pandas matplotlib scikit-learn
15
16 # 6. When finished, deactivate
17 deactivate
```

## 9.2 Cloning a Shared GPU-Optimized TensorFlow Environment

```
1 # 1. Install the clone tool (only once per user)
2 pip3 install --user virtualenv-clone
3
4 # 2. Create a directory to store your environments (if not already created)
5 mkdir -p ~/myEnvs
6
7 # 3. Clone the shared TensorFlow environment into your home directory
8 virtualenv-clone /shared/tensorflow-2.6.2/tf_env ~/myEnvs/tf_env
9
10 # 4. Activate the cloned environment
11 source ~/myEnvs/tf_env/bin/activate
12
13 # 5. Install extra libraries needed for your project
14 pip install pandas scikit-learn matplotlib seaborn tqdm
15
16 # 6. Deactivate when done
17 deactivate
```

## 9.3 Selecting an Environment in a Job Script

### Activate the CPU-only environment:

```
1 # Inside your SLURM job script
2 echo "Activating CPU-only environment"
3 source ~/myEnvs/cpu_env/bin/activate
```

### Activate the cloned GPU TensorFlow environment:

```
1 # Inside your SLURM job script
2 echo "Activating GPU TensorFlow environment"
3 source ~/myEnvs/tf_env/bin/activate
```

## 9.4 Installing Libraries After Environment Activation

Inside the CPU-only environment:

```
1 source ~/myEnvs/cpu_env/bin/activate  
2 pip install jupyter notebook ipykernel  
3 deactivate
```

Inside the cloned GPU environment:

```
1 source ~/myEnvs/tf_env/bin/activate  
2 pip install opencv-python pillow tensorflow  
3 deactivate
```

## 9.5 Checking and Verifying Resources

This subsection shows how to verify what resources are available on the cluster and what resources SLURM has actually allocated to your jobs.

### 9.5.1 Check Available Partitions and Nodes

To see which partitions and nodes exist and their current status:

```
1 # Show all partitions and their status  
2 sinfo
```

For more detailed information about each node:

```
1 # Detailed node information  
2 scontrol show nodes
```

### 9.5.2 Check Your Jobs

To see all of your jobs in the queue:

```
1 # Show jobs for your user  
2 squeue -u $USER
```

To inspect a specific job by its job ID:

```
1 # Show a specific job  
2 squeue -j <JOBID>
```

### 9.5.3 Inspect Resources Allocated by SLURM

To see precisely what resources were allocated to a job:

```
1 # Show full job allocation details  
2 scontrol show job <JOBID>
```

Look for fields such as NumNodes, NumCPUs, Gres, and NodeList.

### 9.5.4 Check Resources Inside a Job Script

Inside a running job, you can query SLURM environment variables to see what was allocated:

```

1 echo "Job ID: $SLURM_JOB_ID"
2 echo "Nodes allocated: $SLURM_NNODES"
3 echo "Node list: $SLURM_NODELIST"
4 echo "CPUs per node: $SLURM_CPUS_ON_NODE"
5 echo "Tasks per node: $SLURM_TASKS_PER_NODE"
6 echo "Job partition: $SLURM_JOB_PARTITION"

```

### 9.5.5 Check GPU Allocation

For GPU jobs, it is important to verify which GPUs are visible:

```

1 # GPUs assigned by SLURM (logical IDs)
2 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
3
4 # List visible GPUs and utilization
5 nvidia-smi

```

### 9.5.6 Check CPU and Memory Usage

You can inspect CPU and memory characteristics on the node:

```

1 # CPU info
2 lscpu
3
4 # Memory info
5 free -h
6
7 # Per-node memory summary (first lines of /proc/meminfo)
8 cat /proc/meminfo | head

```

### 9.5.7 Check Resources Used by Completed Jobs

After a job finishes, accounting information can be retrieved with `sacct`:

```

1 # Accounting info for a finished job
2 sacct -j <JOBID> --format=JobID,JobName,Partition,AllocCPUS,Elapsed,MaxRSS,State

```

This shows how many CPUs were allocated, how long the job ran, peak memory usage (MaxRSS), and the final state (e.g., COMPLETED, FAILED, TIMEOUT).

### 9.5.8 Minimal Resource Debug Job

The following SLURM script does nothing except report which resources it received. It is useful when testing new partitions or debugging allocation issues:

```

1  #!/bin/bash
2  #SBATCH --job-name=resource_debug
3  #SBATCH -N 1
4  #SBATCH -p kimq
5  #SBATCH --gres=gpu:1
6  #SBATCH --time=00:10:00
7
8  echo "Job ID: $SLURM_JOB_ID"
9  echo "Node list: $SLURM_NODELIST"
10 echo "CPUs/node: $SLURM_CPUS_ON_NODE"
11 echo "GPUs visible: $CUDA_VISIBLE_DEVICES"
12
13 echo "===== nvidia-smi ====="
14 nvidia-smi
15
16 echo "===== lscpu ====="
17 lscpu
18
19 echo "===== memory ====="
20 free -h

```

## 9.6 SLURM Job Templates for Different Resource Requirements

This subsection shows full job scripts for common use cases.

### 9.6.1 Single Node, CPU-Only Job

```

1  #!/bin/bash
2  #SBATCH --job-name=cpu_only_job
3  #SBATCH --output=cpu_only_job.out.%j
4  #SBATCH --error=cpu_only_job.err.%j
5  #SBATCH -N 1
6  #SBATCH -p compute # CPU-only partition
7  #SBATCH --time=01:00:00 # 1 hour
8
9  echo "Host: $(hostname)"
10 echo "Using CPU-only environment"
11 source ~/myEnvs/cpu_env/bin/activate
12
13 echo "Running CPU-only Python script"
14 python ~/projects/my_cpu_script.py
15
16 echo "Deactivating environment"
17 deactivate

```

### 9.6.2 Single Node, Single GPU Job

```
1 #!/bin/bash
2 #SBATCH --job-name=single_gpu_job
3 #SBATCH --output=single_gpu_job.out.%j
4 #SBATCH --error=single_gpu_job.err.%j
5 #SBATCH -N 1
6 #SBATCH -p kimq # GPU partition
7 #SBATCH --gres=gpu:1 # 1 GPU
8 #SBATCH --time=02:00:00 # 2 hours
9
10 echo "Host: $(hostname)"
11 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
12
13 echo "Activating GPU TensorFlow environment"
14 source ~/myEnvs/tf_env/bin/activate
15
16 echo "Running single-GPU training script"
17 python ~/projects/train_single_gpu.py
18
19 echo "Deactivating environment"
20 deactivate
```

### 9.6.3 Single Node, Multiple GPUs (e.g., 4 GPUs)

```
1 #!/bin/bash
2 #SBATCH --job-name=multi_gpu_job
3 #SBATCH --output=multi_gpu_job.out.%j
4 #SBATCH --error=multi_gpu_job.err.%j
5 #SBATCH -N 1
6 #SBATCH -p kimq # GPU partition
7 #SBATCH --gres=gpu:4 # 4 GPUs on the same node
8 #SBATCH --time=04:00:00 # 4 hours
9
10 echo "Host: $(hostname)"
11 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
12
13 echo "Activating GPU TensorFlow environment"
14 source ~/myEnvs/tf_env/bin/activate
15
16 echo "Running multi-GPU training script (e.g., data parallel)"
17 python ~/projects/train_multi_gpu.py --num-gpus 4
18
19 echo "Deactivating environment"
20 deactivate
```

#### 9.6.4 Multiple Nodes, CPU-Only (e.g., MPI or Distributed CPU Job)

```
1  #!/bin/bash
2  #SBATCH --job-name=cpu_multi_node_job
3  #SBATCH --output=cpu_multi_node_job.out.%j
4  #SBATCH --error=cpu_multi_node_job.err.%j
5  #SBATCH -N 2 # 2 nodes
6  #SBATCH -p compute # CPU-only partition
7  #SBATCH --time=03:00:00 # 3 hours
8
9  echo "SLURM_NODELIST: $SLURM_NODELIST"
10 echo "SLURM_NNODES: $SLURM_NNODES"
11
12 echo "Activating CPU-only environment"
13 source ~/myEnvs/cpu_env/bin/activate
14
15 # Example: launch an MPI or distributed job
16 # (replace with your actual launcher, e.g., srun mpirun, torchrun, etc.)
17 echo "Running distributed CPU job"
18 srun python ~/projects/distributed_cpu_job.py
19
20 echo "Deactivating environment"
21 deactivate
```

#### 9.6.5 Multiple Nodes, Multiple GPUs (Distributed GPU Training)

```
1  #!/bin/bash
2  #SBATCH --job-name=multi_node_multi_gpu
3  #SBATCH --output=multi_node_multi_gpu.out.%j
4  #SBATCH --error=multi_node_multi_gpu.err.%j
5  #SBATCH -N 2 # 2 nodes
6  #SBATCH -p kimg # GPU partition
7  #SBATCH --gres=gpu:4 # 4 GPUs per node (total 8 GPUs)
8  #SBATCH --time=06:00:00 # 6 hours
9
10 echo "SLURM_NODELIST: $SLURM_NODELIST"
11 echo "SLURM_NNODES: $SLURM_NNODES"
12 echo "CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"
13
14 echo "Activating GPU TensorFlow environment"
15 source ~/myEnvs/tf_env/bin/activate
16
17 # Example using srun to launch a distributed training job
18 # Replace with your actual distributed launcher (e.g., torchrun, horovodrun)
19 echo "Running multi-node, multi-GPU training"
20 srun python ~/projects/train_distributed_gpu.py \
21     --nodes $SLURM_NNODES \
22     --gpus-per-node 4
23
24 echo "Deactivating environment"
25 deactivate
```

## 9.7 Quick Mapping: Which Example to Use?

- **Create new CPU environment:** Listing in Section [9.1](#).
- **Clone shared GPU TF environment:** Listing in Section [9.2](#).
- **CPU-only job:** Section [9.6.1](#).
- **Single-GPU job:** Section [9.6.2](#).
- **Single node, multiple GPUs:** Section [9.6.3](#).
- **Multi-node CPU job:** Section [9.6.4](#).
- **Multi-node, multi-GPU job:** Section [9.6.5](#).

## 10 Summary

In this tutorial, you have seen:

- How to write and submit a basic SLURM job script that runs a TensorFlow program.
- How to interpret key `#SBATCH` directives: number of nodes, partition, GPU requests, and time limits.
- Why shared environments on HPC clusters are read-only and how to safely extend them by cloning to your home directory.
- How to create your own Python virtual environments from scratch and how to connect them to SLURM job scripts.

This workflow lets you run fast, GPU-enabled jobs using cluster-optimized software while keeping your custom Python packages isolated and safe.

# Index

-gres, 5  
-time, 5  
-N, 4  
-p, 5  
#SBATCH directives, 3

Conda environment  
    IR, 11

CRADLE, 3

CUDA, 5

cuDNN, 5

environment  
    shared, 3

GPU  
    partition, 5  
    request, 5

GPU workload, 3

GRES, 5

HPC cluster, 1

Micromamba, 10

node  
    compute, 4

pandas, 7

partition, 5

Python environment, 8

sacct, 17

scheduler  
    directives, 4

scp, 11

SLURM, 1  
    batch script, 3  
    environment variables, 16  
    resources, 16

ssh, 3

TensorFlow, 3  
    shared installation, 5

time limit, 5

university cluster, 3

venv, 8  
virtual environment, 3  
virtualenv-clone, 6  
VPN, 3